

Learning Sentiment-Specific Word Embedding (SSWE) for Twitter Sentiment Classification

Chang Shu

04/02/2020

Introduction

- Sentiment classification
 - Label a text's sentiment polarity as “positive”, “negative”, and “neutral”
- Why sentiment classification?
 - Commerce
 - Consumer's attitude towards commercial product and competitors
 - Reveal confidence of the market to predict stock price
 - AI
 - Help AI understand human feelings to build better chatbot
 - Politics
 - public opinion to policy announcements and campaign messages
- Why twitter?
 - Larger
 - Prompt
 - General (compared with blogpost, forum)

How?

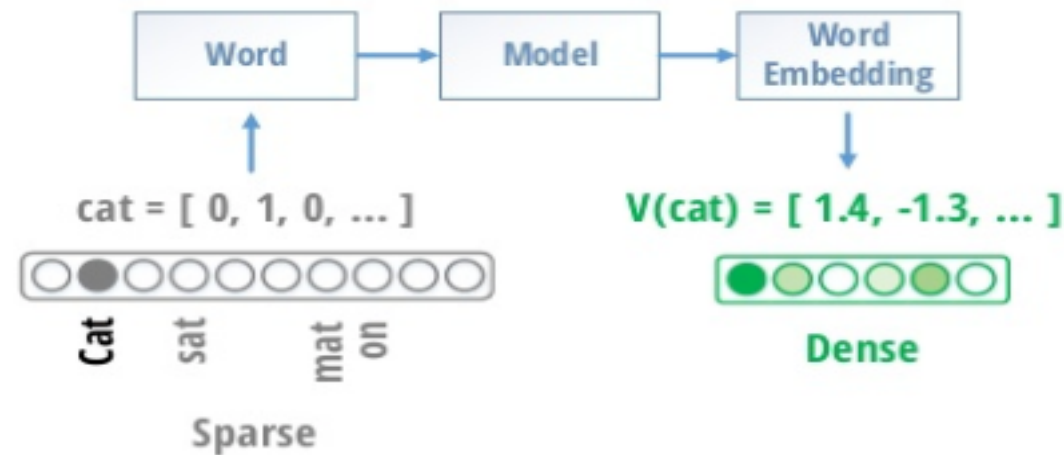
- The majority of existing approaches employ machine learning algorithms to build Tweets sentiment classifier.
 - Sentiment Lexicon
 - Feature Learning method

Previous work: sentiment lexicon

- Sentiment lexicons are list of words with associations to positive and negative sentiments (*Mohammad et al.,2013*)
- *Mohammad et al.(2013)* build the top performance Tweets sentiment classification using sentiment lexicons
 - calculate sentiment association score of each term
 - Hand crafted feature: # of positive terms in tweets
- Cons: labor-intensive, need to less depend on extensive feature engineering

Previous work: syntactic based feature learning method

- Word embedding
 - Each word was projected from sparse representation (one-hot) to a dense, low-dimensional and real valued vector
 - C&W model (*Collobert et al., 2011*), Word2Vec (*Mikolov et al., 2013*)



C&W Model


Original tweet

cat	chills	on	a	mat
-----	--------	----	---	-----

Corrupted tweet

cat	chills	job	a	mat
-----	--------	-----	---	-----

- Method
 - Replace the center word with a random word and create a corrupted tweet
 - Calculate language model score for each original tweet and its corrupted tweet
- The training objective
 - is that the original n gram is expected to obtain a higher language model score than the corrupted n gram by a margin of 1
- Cons: only model the syntactic context but ignore the sentiment information
 - “good” and “bad” are mapped into close vector
- Calls for combining sentiment context as well



Algorithms comparison

Model	Syntactic	Sentiment
C&W	YES	NO
SSWE(h)	NO	YES
SSWE(r)	NO	YES
SSWE(u)	YES	YES

* Sentiment-Specific
Word Embedding
(SSWE)

The authors' approaches, will be introduced in the coming slides

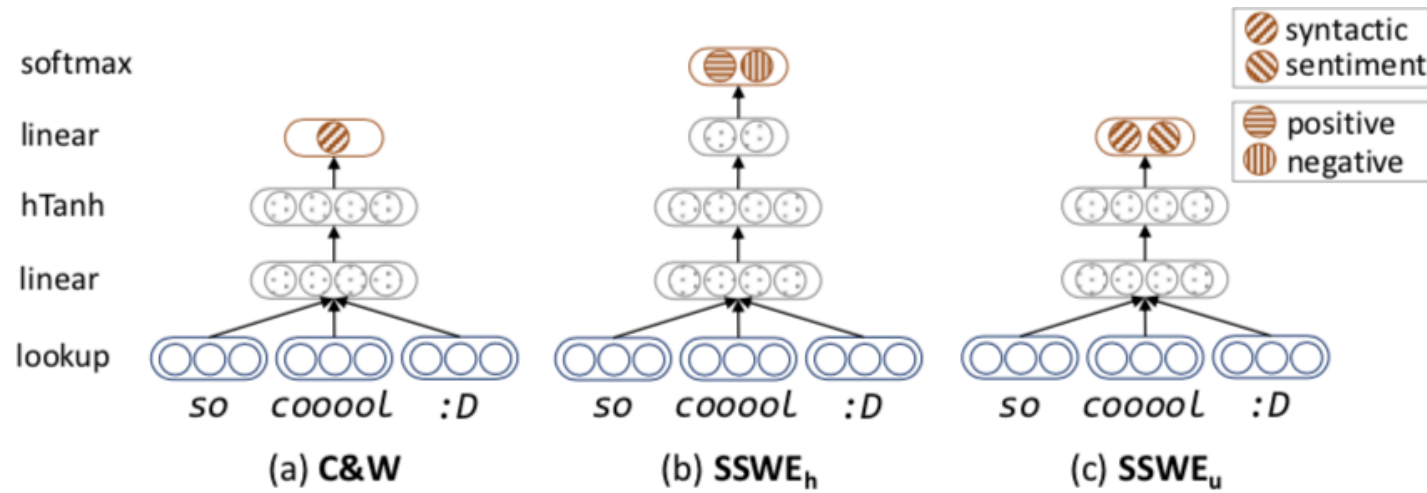
SSWE(h)

Original tweet

cat	chills	on	a	mat
-----	--------	----	---	-----

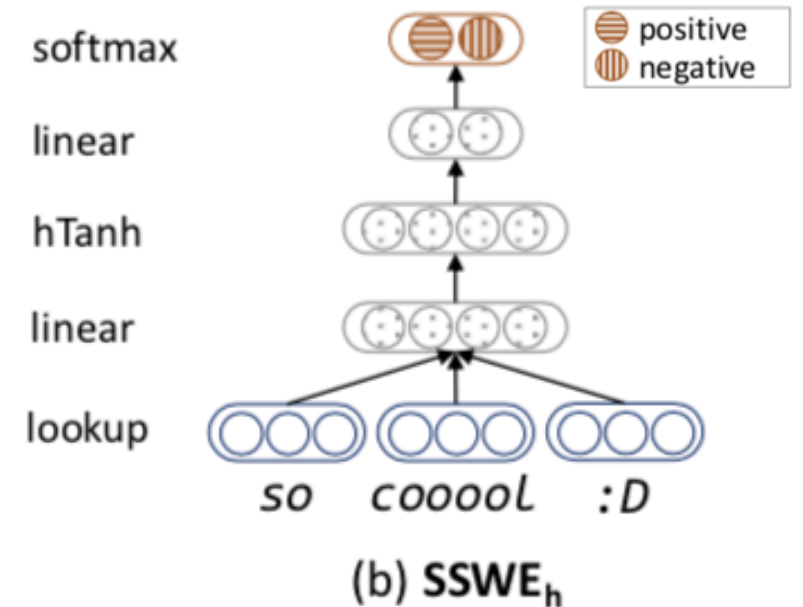
lookup

$[1,0,0,0,0,\dots]$ $[0,0,0,0,1,\dots]$




SSWE(h)

- Learn sentiment specific information
 - Predict sentiment classification for each tweet
 - Ground truth: [1, 0] as positive, [0, 1] as negative
 - Sample output: [0.7, 0.3]
- Process
 1. Lookup corresponding word embeddings
 2. Linear projection to lower dimension
 3. Hyperbolic tangent to constraint the range
 4. Linear projection to # classes (2)
 5. Softmax for prediction



SSWE(*r*): *r* relaxed version of SSWE(*h*)

- Relax the hard constraint of SSWE(*h*)
 - SSWE(*h*) define [1,0] as positive sentiment and [0,1] as negative
 - After relax the constraint, SSWE(*r*)'s output [*a*,*b*], if *a* > *b*, then positive, and vice versa
 - Borrow the first 4 layer from SSWE(*h*) and remove softmax layer




SSWE(**u**): **u**nified sentiment and syntactic context

Original tweet	cat	chills	on	a	mat
Corrupted tweet	cat	chills	job	a	mat


$$loss_u(t, t^r) = \alpha \cdot loss_{cw}(t, t^r) + (1 - \alpha) \cdot loss_{us}(t, t^r)$$

- “ t ” is the original tweet and “ t^r ” is corrupted tweet
- $loss_{cw}(t, t^r)$ is the syntactic loss
- $loss_{us}(t, t^r)$ is the sentiment loss
- α is the hyper-parameter that weights two parts




Experiment

- Word embedding learning dataset
 - 10M Tweets (04/01/2013-04/30/2013), 5M positive and 5M negative, detected by emoticons.
 - “We use the emoticons selected by Hu et al. (2013). The positive emoticons are :) :) :-) :D =), and the negative emoticons are :(: (:- (“
- Experiment dataset:
 - Twitter sentiment classification benchmark dataset: SemEval2013



	Positive	Negative	Neutral	Total
Train	2,642	994	3,436	7,072
Dev	408	219	493	1,120
Test	1,570	601	1,639	3,810



Result

Embedding	Macro-F1
C&W	74.89
Word2vec	73.21
ReEmb(C&W)	75.87
ReEmb(w2v)	75.21
WVSA	77.04
SSWE _h	81.33
SSWE _r	80.45
SSWE _u	83.70

Table: Macro-F1 on positive/negative classification of tweets with different word embeddings.

*Evaluation Metric: Macro-F1

Result




Figure 2: Macro-F1 of $SSWE_u$ on the development set of SemEval 2013 with different α .

Weakness

- The embedding training dataset is biased by its twitter crawling criteria: using emoticon to determine sentiment polarity
- Human language expression is more than that (sarcasm):
 - E.g., I love doing my homework :)
 - This could mean: 1) I really love doing homework or 2) Nah, I'm just kidding
- The sarcasm, which usually comes with positive emoticon, will be labeled as positive sentiment

Future work

- Better ways to replace the twitter sentiment labeling method instead of using emoticon:
 - E.g., use online product review, classify sentiment polarity by the stars given

Thank you!

Back up

- Min, Max, Average of “embeddings”

$$z(tw) = [z_{max}(tw), z_{min}(tw), z_{average}(tw)]$$

Back up

- SSWE(u) loss function

$$loss_u(t, t^r) = \alpha \cdot loss_{cw}(t, t^r) + (1 - \alpha) \cdot loss_{us}(t, t^r)$$

$$loss_{cw}(t, t^r) = \max(0, 1 - f^{cw}(t) + f^{cw}(t^r))$$

$$loss_{us}(t, t^r) = \max(0, 1 - \delta_s(t) \mathbf{f}_1^u(t) + \delta_s(t) \mathbf{f}_1^u(t^r))$$

Weakness

- Authors do not provide performance comparison between SSWE(h), SSWE(r) and SSWE(u)
 - Only SSWE(u) performance is compared with other state-of-the-art models
 - If SSWE(h)/SSWE(r) has close/better performance than SSWE(u), then there's no need to use syntax information